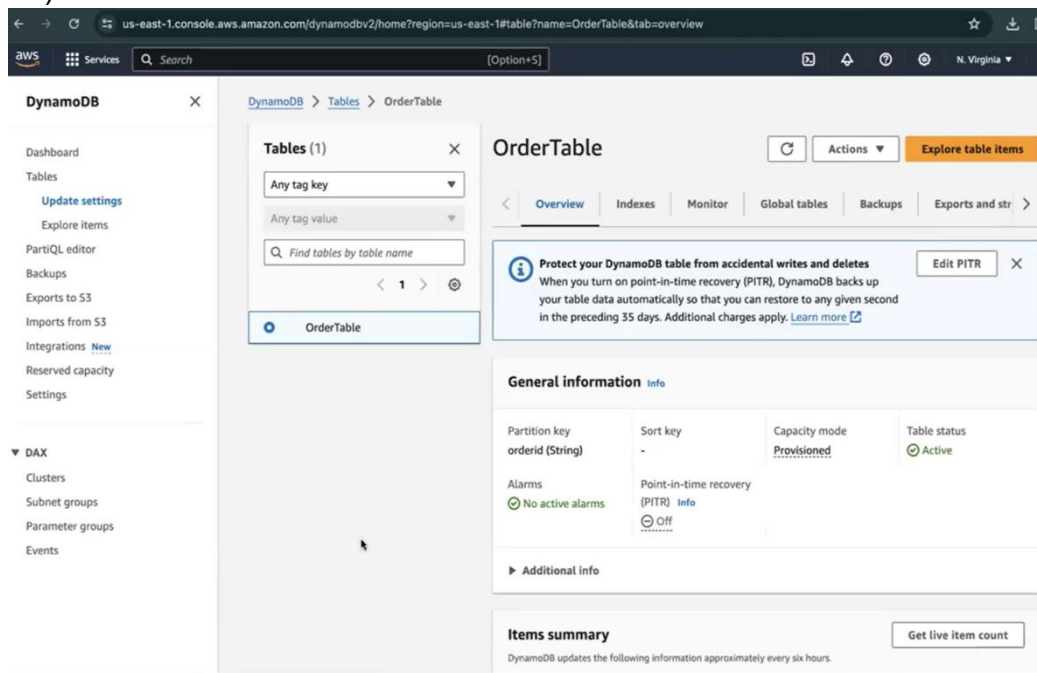
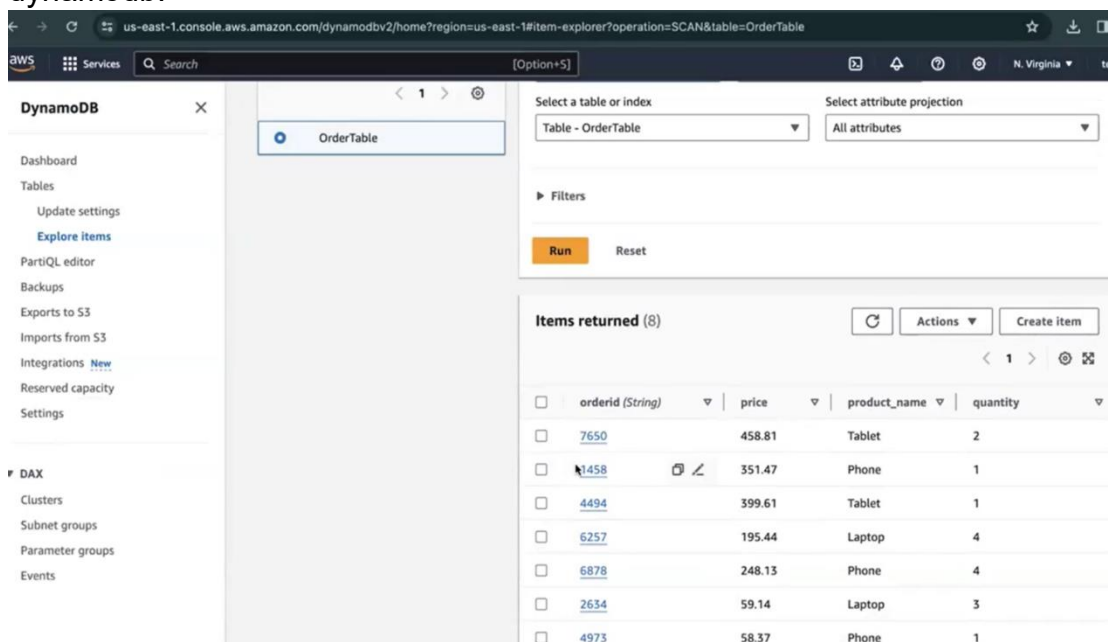


Sales Data Projection

1. this was a sales data streaming project for processing the data from Dynamo db to destination S3 target path.
2. there was some source from where records(orderid, product_name, quantity, price) were being written into the nosql dynamodb(which is a key value based db).



3. we created a dynamodb table where we targeted particular partition key(orderid) and the entire data(orderid, product_name, quantity, price) used to get ingested as a value there. That's how we were publishing a record in dynamodb.



- next our task was to capture changes happening in table. Records were getting inserted, updated/deleted to table in real time. CDC is whatever changes are happening in real time we capture those changes. That's what CDC is. So we needed to enable dynamodb streams from 'Exports & Streams' option to perform CDC. As we were able to consume the CDC changes, hence created CDC pipeline.
- then we set up a kinesis stream which is a queuing mechanism to hold our data. Kinesis stream has a concept of shards just like publishing data to partitions. Data here arrives at kinesis stream in different shard.

The screenshot shows the Amazon Kinesis console interface. The left sidebar contains navigation links for Dashboard, Data streams, Data Firehose, Managed Apache Flink, and Resources. The main content area displays the 'Data stream summary' for the 'kinesis-sales-order' stream. The summary includes the following details:

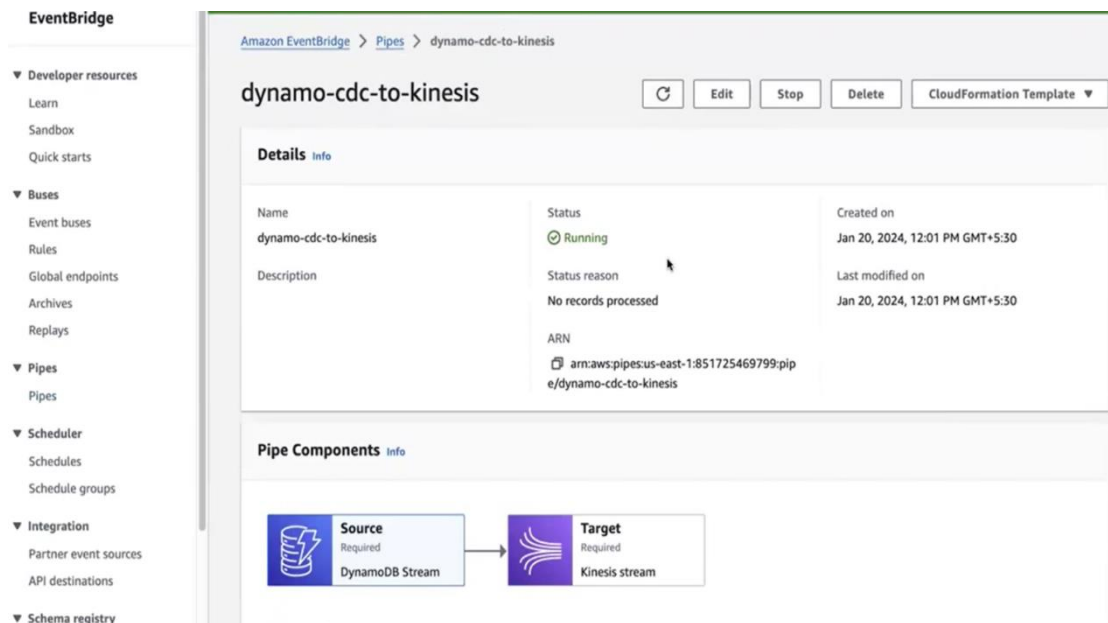
Status	Capacity mode	ARN	Creation time
Active	On-demand	arn:aws:kinesis:us-east-1:851725469799:stream/kinesis-sales-order	January 20, 2024 at 11:53 GMT+5:30

Below the summary, the 'Data viewer' tab is selected, showing a list of records for the selected shard (shardId-000000000002). The records are displayed in a table with columns for Partition key and Data. The first record is an INSERT event for a product named 'Charger'.

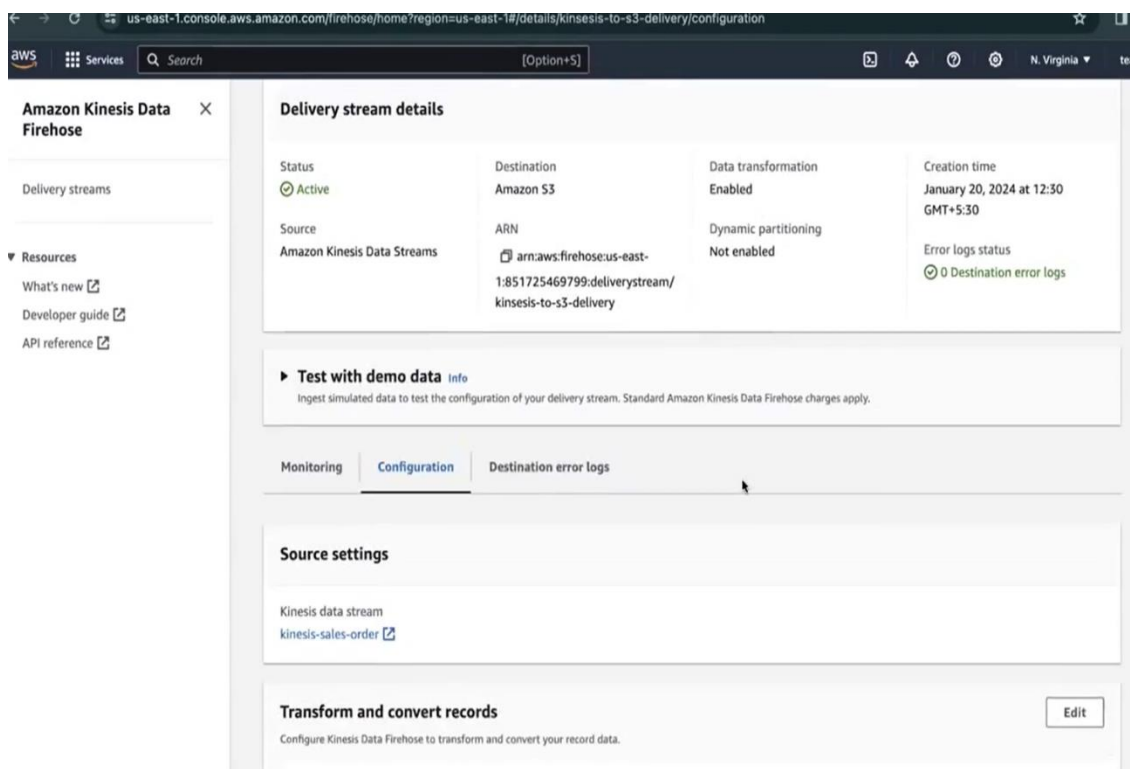
The screenshot shows the 'Record data' modal in the Amazon Kinesis console. The modal displays the sequence number and shard ID of the record. The 'Raw data' tab is selected, showing the raw JSON data of the record. The data is as follows:

```
{
  "eventID": "f3d2479e0c6ed4762e99a377f4a143b6",
  "eventName": "INSERT",
  "eventVersion": "1.1",
  "eventSource": "aws:dynamodb",
  "awsRegion": "us-east-1",
  "dynamodb": {
    "ApproximateCreationDateTime": 1705732562,
    "Keys": {
      "orderid": {
        "S": "3714"
      }
    },
    "Newimage": {
      "quantity": {
        "N": "5"
      },
      "orderid": {
        "S": "3714"
      },
      "price": {
        "N": "190.07"
      },
      "product_name": {
        "S": "Charger"
      },
      "SequenceNumber": "3510000000077909614008",
      "SizeBytes": 60,
      "StreamViewType": "NEW_IMAGE"
    },
    "eventSourceARN": "arn:aws:dynamodb:us-east-1:851725469799:table/OrderTable/stream/2024-01-20T06:21:09.489"
  }
}
```

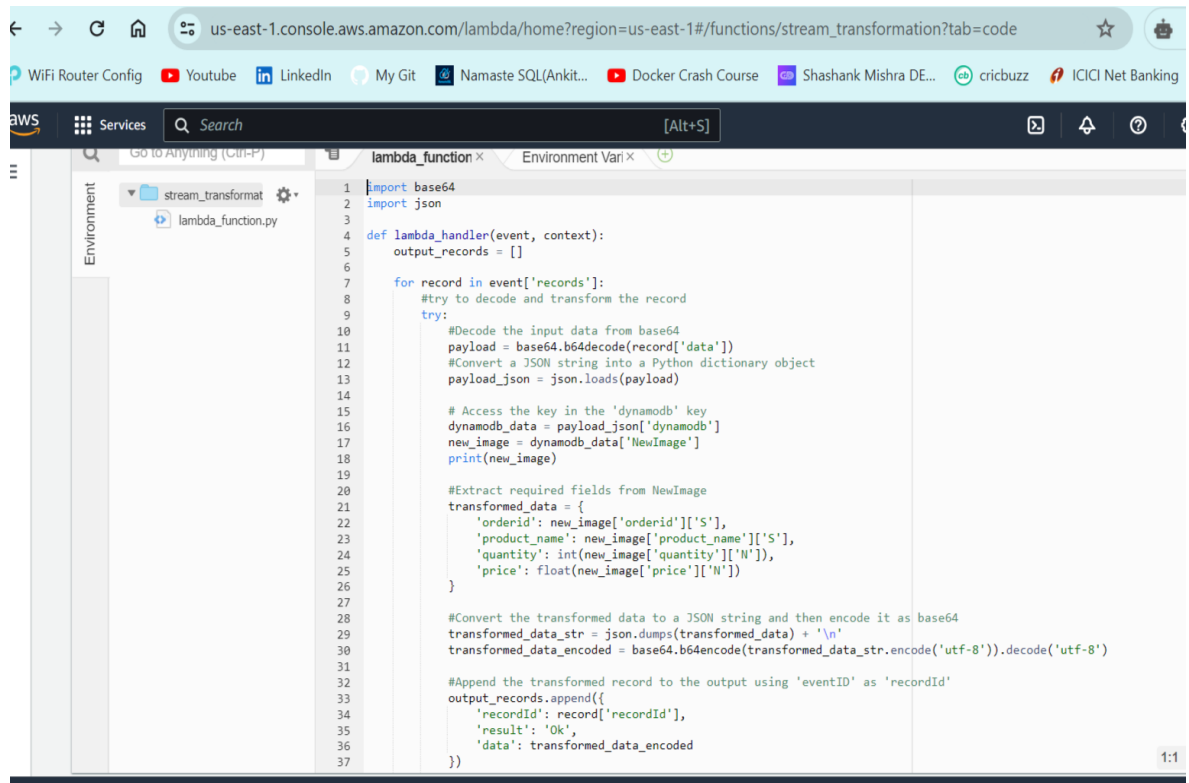
6. then we had a event bridge pipe in between dynamodb stream and kinesis stream. We set a batch size consisting of no of message to be sent to any kinesis stream shard. Set up required IAM role to eventbridge pipe to access dynamo stream & kinesis stream. Using event bridge we used to flow data to kinesis stream.



7. next we created a firehose stream which is a delivery stream. It processed data(like ingest, transform) and delivered streaming data to destinations(like data lakes, data warehouses). Here source is kinesis stream and destination is s3.



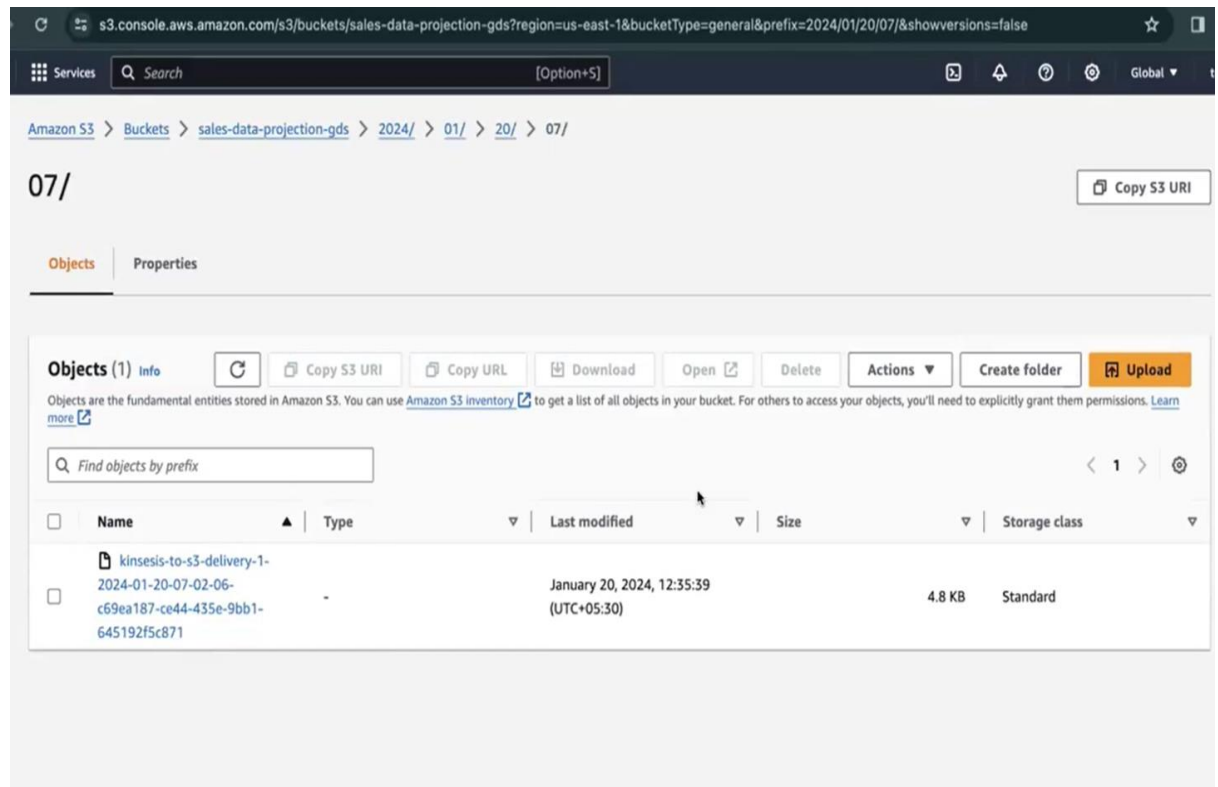
8. using delivery stream we were batching the records and transforming those mini batch records with the help of lambda function.



The screenshot shows the AWS Lambda console for a function named 'stream_transform'. The code is written in Python and is designed to process records from a delivery stream. It imports 'base64' and 'json' modules. The 'lambda_handler' function takes an event and context as input. It iterates over the 'records' in the event, decoding each record from base64, loading it as JSON, and then extracting specific fields ('orderid', 'product_name', 'quantity', 'price') to create a 'transformed_data' dictionary. This data is then converted back to a JSON string, encoded to base64, and appended to the 'output_records' list. The function finally returns the 'output_records'.

```
1 import base64
2 import json
3
4 def lambda_handler(event, context):
5     output_records = []
6
7     for record in event['records']:
8         #try to decode and transform the record
9         try:
10             #Decode the input data from base64
11             payload = base64.b64decode(record['data'])
12             #Convert a JSON string into a Python dictionary object
13             payload_json = json.loads(payload)
14
15             # Access the key in the 'dynamodb' key
16             dynamodb_data = payload_json['dynamodb']
17             new_image = dynamodb_data['NewImage']
18             print(new_image)
19
20             #Extract required fields from NewImage
21             transformed_data = {
22                 'orderid': new_image['orderid']['S'],
23                 'product_name': new_image['product_name']['S'],
24                 'quantity': int(new_image['quantity']['N']),
25                 'price': float(new_image['price']['N'])
26             }
27
28             #Convert the transformed data to a JSON string and then encode it as base64
29             transformed_data_str = json.dumps(transformed_data) + '\n'
30             transformed_data_encoded = base64.b64encode(transformed_data_str.encode('utf-8')).decode('utf-8')
31
32             #Append the transformed record to the output using 'eventId' as 'recordId'
33             output_records.append({
34                 'recordId': record['recordId'],
35                 'result': 'Ok',
36                 'data': transformed_data_encoded
37             })
38
39     return {'records': output_records}
```

9. Once this is done, then delivering the transformed batch records to S3.



10. Lastly, we created crawler over s3 target path which resulting in catalog table creation. Also, needed to add classifier for json format(i.e. \$.columns).

The screenshot shows the AWS Glue console interface. On the left is a navigation sidebar with options like 'Getting started', 'Data Catalog tables', 'Data Catalog', 'Data connections', 'Workflows (orchestration)', 'Data Catalog', 'Databases', 'Tables', 'Stream schema registries', 'Schemas', 'Connections', 'Crawlers', 'Classifiers', 'Catalog settings', 'Data Integration and ETL', and 'Legacy pages'. The main content area displays the configuration for the 'sales_data_crawler'. A green banner at the top states 'Crawler successfully starting' and 'The following crawler is now starting: sales_data_crawler'. Below this, the crawler's properties are listed: Name (sales_data_crawler), IAM role (glue-role), Database (sales-data-catalog), State (READY), Description, Security configuration, Lake Formation configuration, and Table prefix (projection_). The 'Advanced settings' section is collapsed. At the bottom, the 'Crawler runs' tab is active, showing a single run from January 20, 2024, at 07:13:30, which is currently 'Running'.

11. using athena query we then did analysis

The screenshot shows the AWS Athena console. On the left, a table named 'projection_sales_data_projection_gds' is listed with columns: orderid (string), product_name (string), quantity (int), price (double), partition_0 (string (Partitioned)), partition_1 (string (Partitioned)), partition_2 (string (Partitioned)), and partition_3 (string (Partitioned)). The main area displays the results of a query. The query status is 'Completed' with a time in queue of 103 ms, a run time of 1.306 sec, and 8.89 KB of data scanned. The results are shown in a table with 10 rows. The columns in the results table are: #, orderid, product_name, quantity, price, partition_0, partition_1, and partition_2. The data shows various products like Phone, Charger, Tablet, Laptop, and Headphones with their respective quantities and prices across different partitions.

#	orderid	product_name	quantity	price	partition_0	partition_1	partition_2
1	7864	Phone	4	175.66	2024	01	20
2	9317	Charger	2	458.18	2024	01	20
3	6468	Charger	3	243.42	2024	01	20
4	7341	Tablet	1	230.83	2024	01	20
5	3958	Charger	3	476.51	2024	01	20
6	6324	Laptop	1	31.55	2024	01	20
7	3238	Headphones	5	434.68	2024	01	20
8	8007	Phone	2	36.83	2024	01	20
9	8033	Phone	3	166.54	2024	01	20